

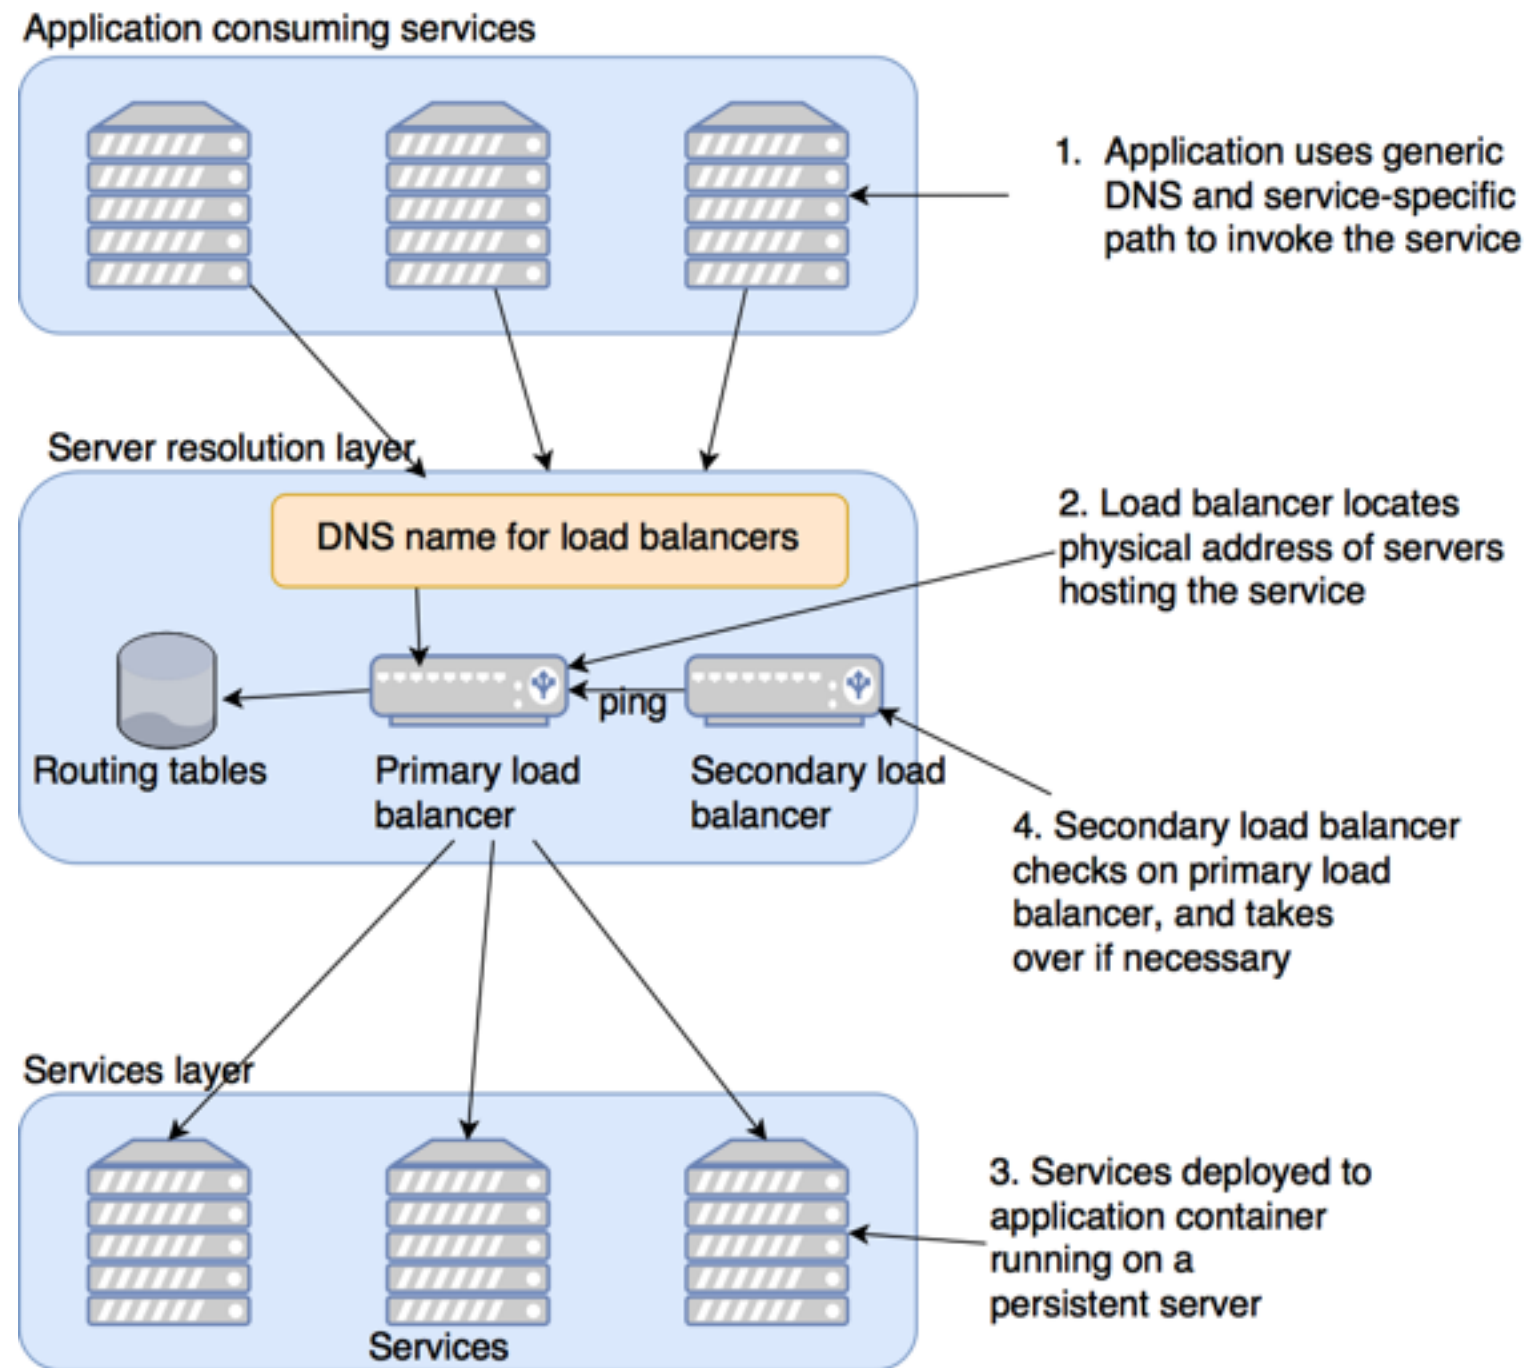
Service Discovery

Chapter Content

- 1. Traditional service location solution model
- 2. Service discovery in the cloud
- 3. Client-side load balancing
- 4. Service discovery using Eureka

- Service discovery is critical to cloud-based microservice for two reasons:
 - It offers the ability to quickly horizontally scale up and down the number of service instances running in an environment
 - It helps increase application resiliency. When an instance fails, it can remove it from the available instance list

1. Traditional service location solution model



- The above architecture works well on-premise, but not for cloud-based microservice applications due to:
 - **Single-point failure** - load balancer is a single point of failure for the entire infrastructure
 - **Limited horizontal scalability** - Hardware constraint
 - **Statically managed** - not designed for rapid registration and de-registration of services
 - **Complex** - service consumer requests have to have their requests mapped to the physical services. This is usually done manually.

2. Service discovery in the cloud

- Four concepts around service discovery architecture:
 - Service registration
 - Client lookup of service address
 - Information sharing
 - Health monitoring



Client application



Service discovery layer

1. A services location can be looked up by a logical name from the service discovery agent.



Service discovery node 1



Service discovery node 2



Service discovery node 3

3. Service discovery nodes share service instance health information among each other.



Service instances

2. When a service comes online it registers its IP address with a service discovery agent.



Service



heart beat



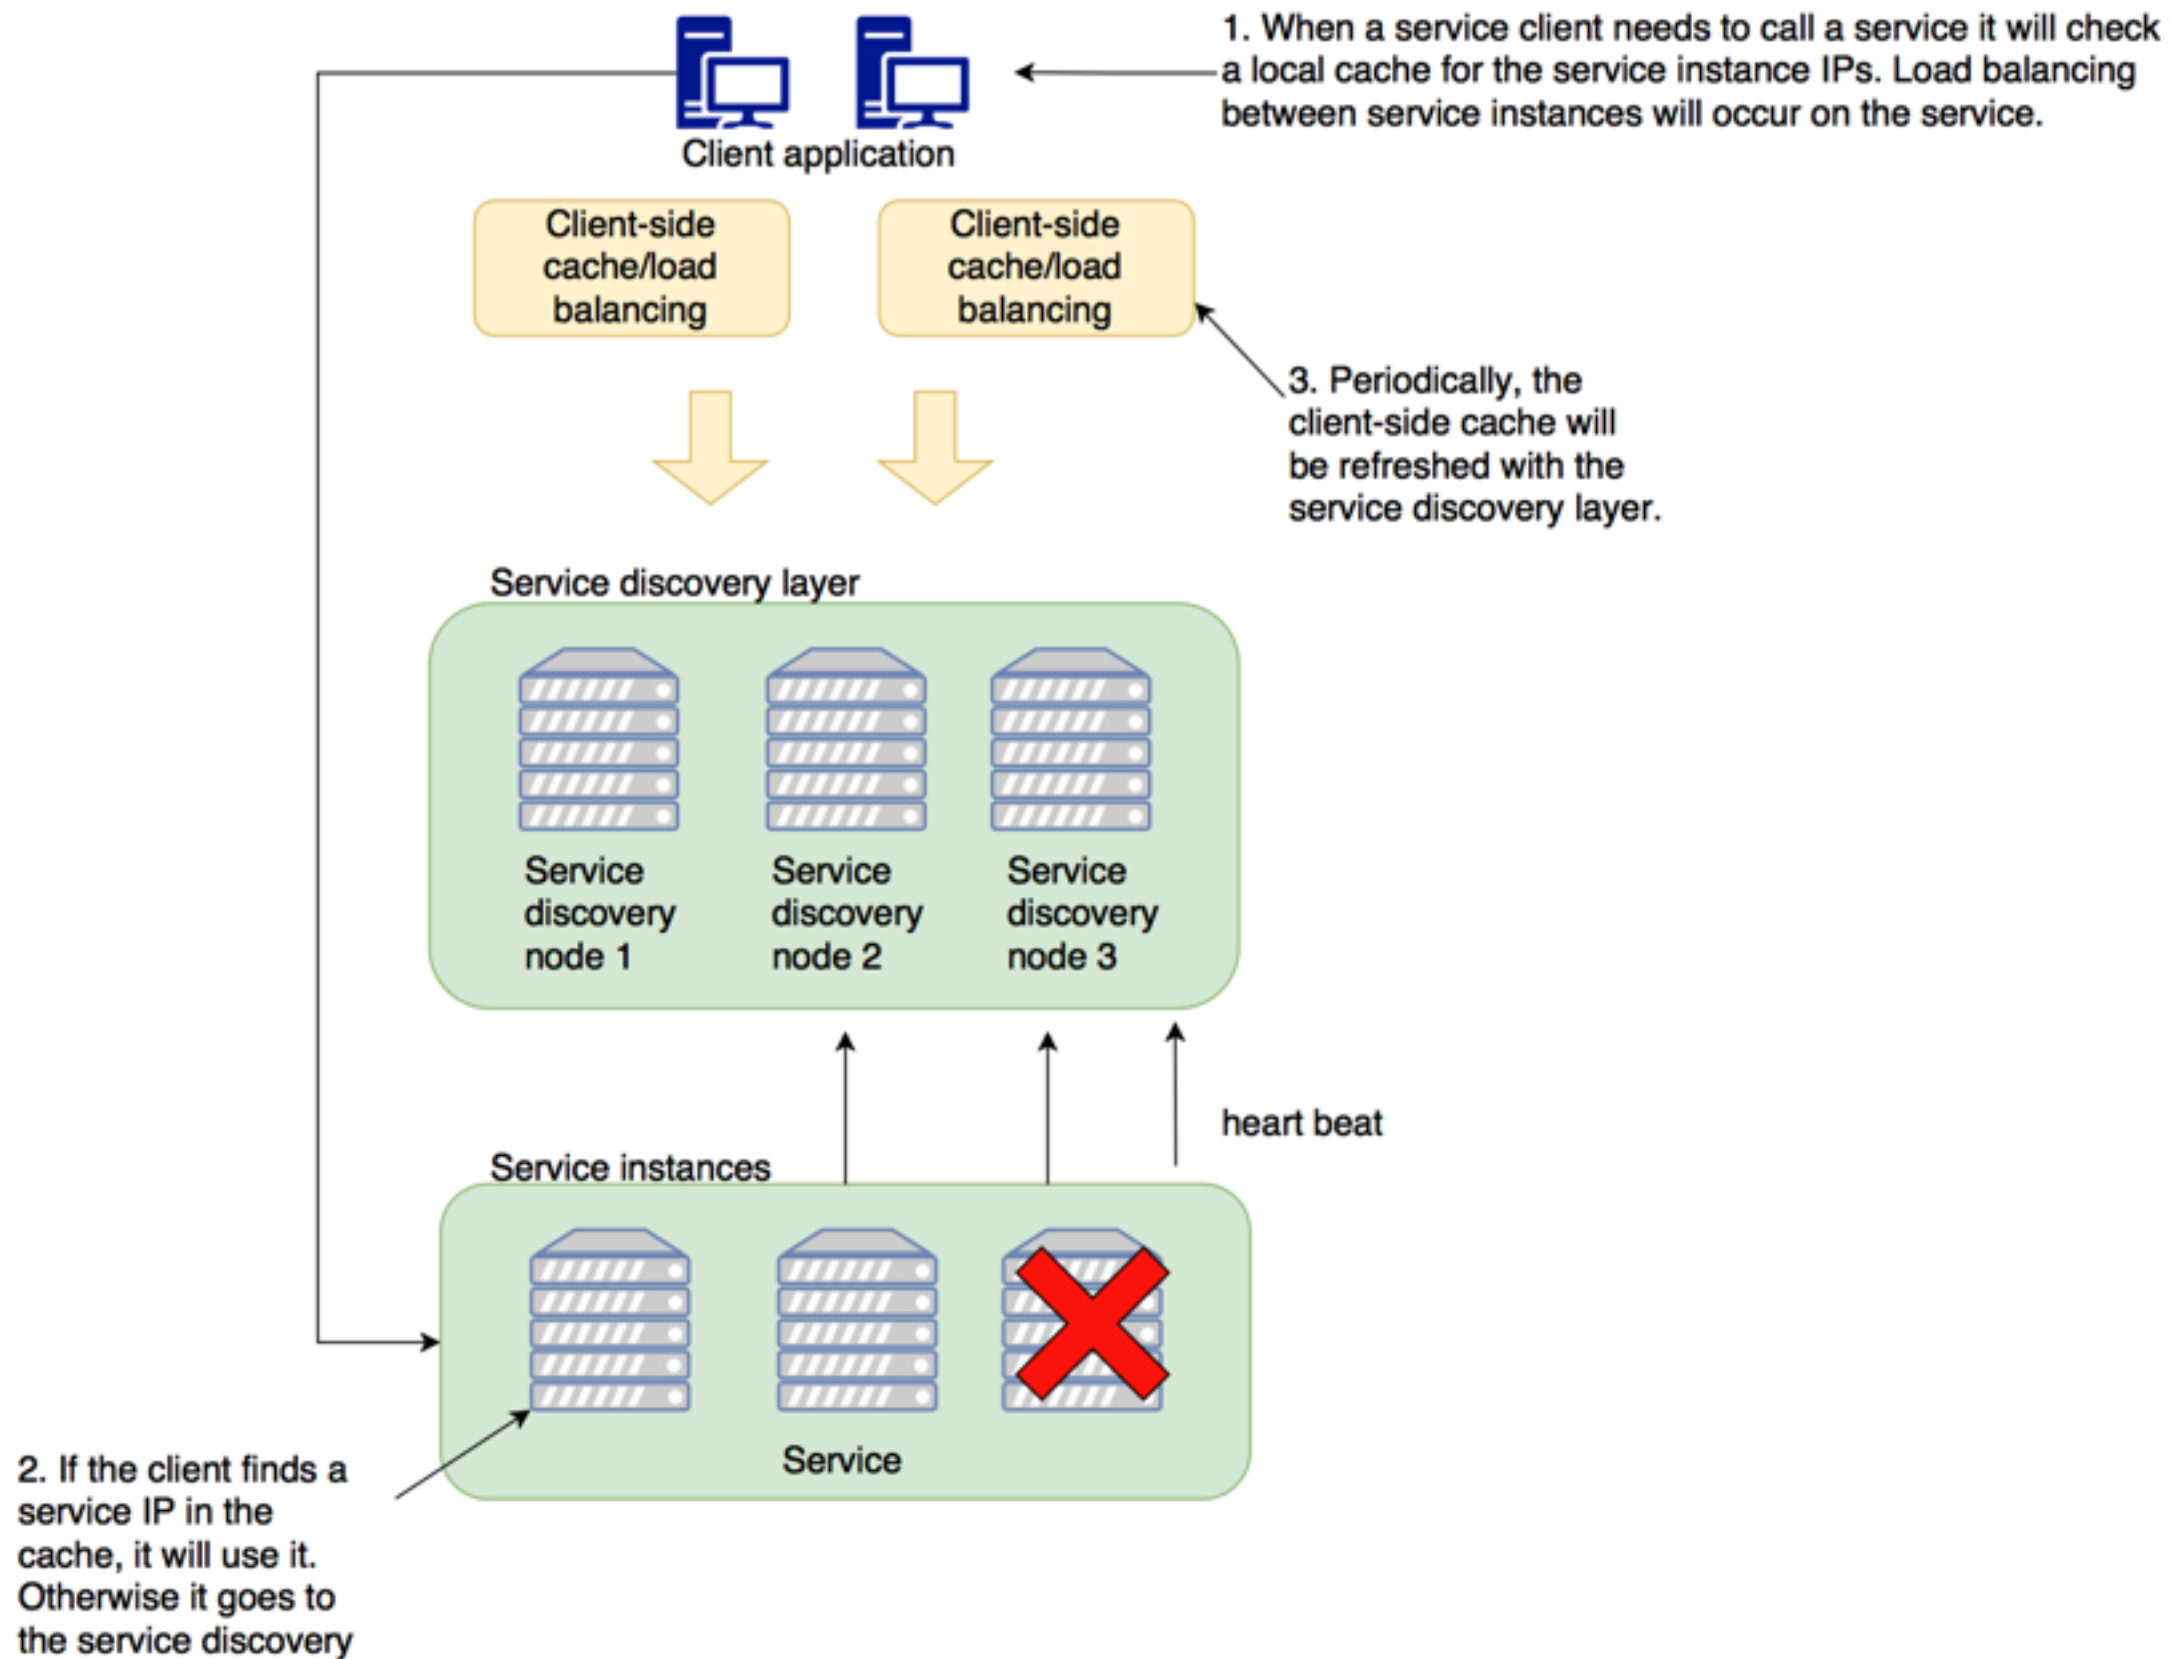
4. Services send a heartbeat to the service discovery agent. If a service dies, the service discovery layer removes the IP of the "dead" instance.



- Discovery agent instances are usually unique and don't have a **load balancer** that sits in front of them
- As service instances start up, they'll register their physical location, path, and port that they can be accessed by with one or more service discovery instances
- While each instance of service will have a **unique IP address** and port, each service instance that comes up will register under the **same service ID**
- A service ID is a key that uniquely identifies a group of the same service instances
- A service will usually **only** register with **one** service discovery service instance. Most service discovery implementations use a peer-to-peer model of data propagation where the data around each service instance is communicated to all the other nodes in the cluster
- Each service instance will push to or have pulled from its status by discovery agent. Any services failing to return a good health check will be removed from the pool of available service instances

3. Client-side load balancing

- There are multiple ways to “discover” a service:
 - A client can rely solely on the service discovery engine to resolve service locations each time a service is called. With this approach, the service discovery engine will be invoked every time a call to a registered microservice instance is made. This approach is brittle as it completely depends on the service discovery engine
 - A more robust approach is to use **client-side load balancing**



4. Service discovery using Eureka

